

Fairness in multimedia analysis

Accuracy-compatible Fairness Computing

Yi Zhang


Co-Contributors: Jitao Sang, Kexin Wang, Yuxuan Yi

Beijing Jiaotong University

yi.zhang@bjtu.edu.cn



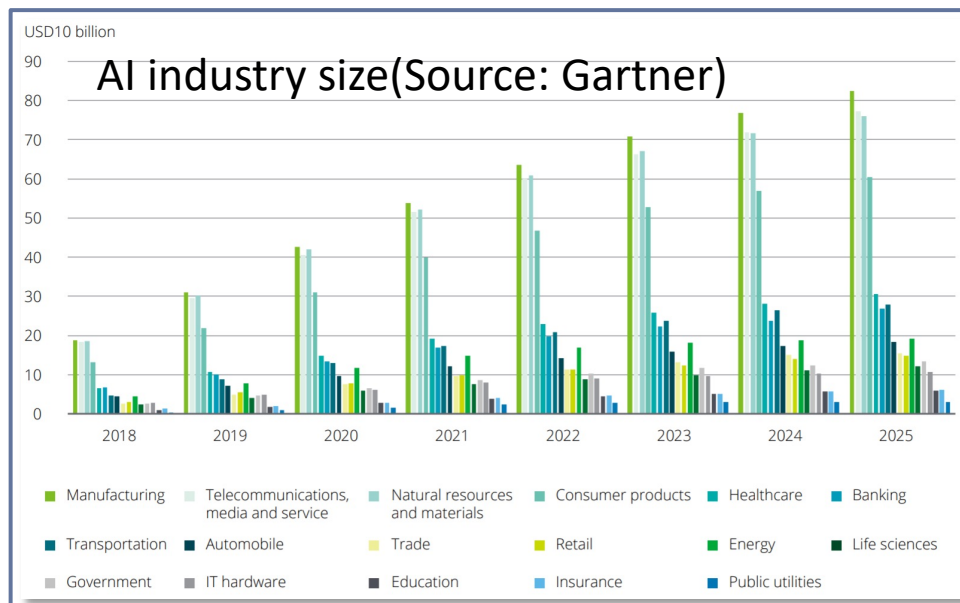
Progress for AI and challenge to fairness



source: <https://www.dsiac.org/resources/journals/dsiac/winter-2017-volume-4-number-1/real-time-situ-intelligent-video-analytics>

Rank	Model	EM	F1
	Human Performance Stanford University (Rajpurkar et al. '16)	82.304	91.221
1	BERT (ensemble) Google A.I.	87.433	93.160
2	BERT (single model) Google A.I.	85.083	91.835
2	nlnet (ensemble) Microsoft Research Asia	85.356	91.202
2	nlnet (ensemble)	85.954	91.677

AI has achieved great success in CV and NLP



Challenge to fairness


What is fairness problem?

Accuracy is not the only one criterion, AI needs criteria in fairness

- ❑ The same algorithm produces **result of discrimination for sensitive people** in decision-making:
 - (1) Different decision-making accuracy for different groups of people;
 - (2) In decision-making, algorithm uses population information that are not related to the task.



Dark face cannot be detected by the face detection model



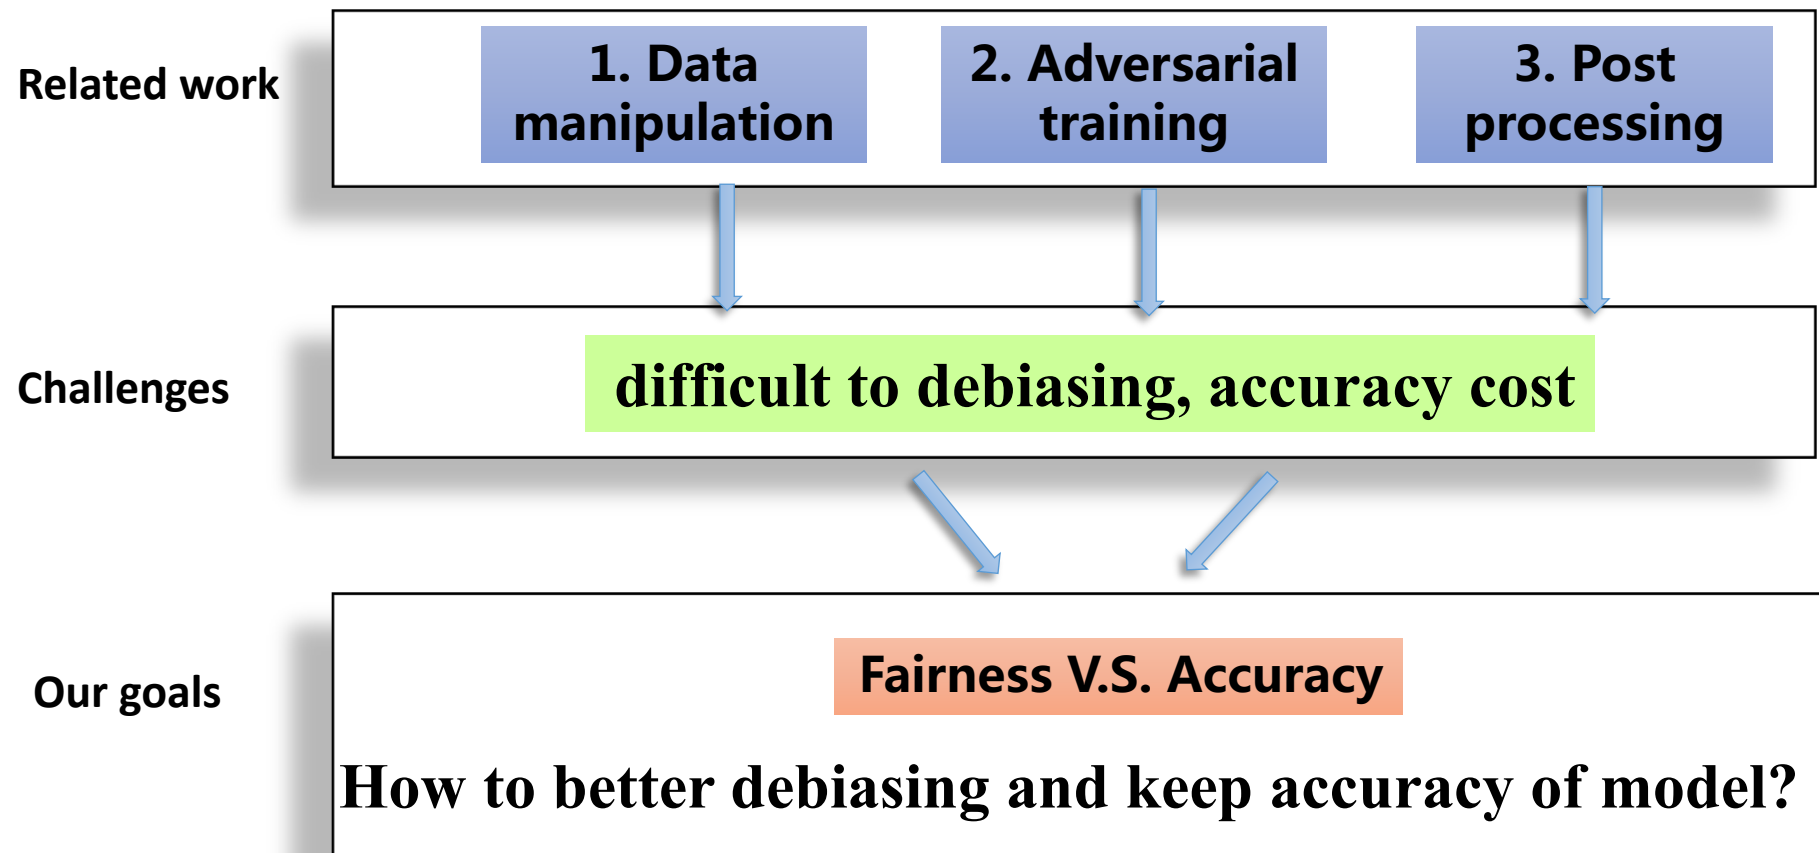
Accuracy gap between darker females and lighter males

DYLAN FUGETT	BERNARD PARKER
Prior Offense 1 attempted burglary	Prior Offense 1 resisting arrest without violence
Subsequent Offenses 3 drug possessions	Subsequent Offenses None
LOW RISK 3	HIGH RISK 10

Judicial algorithm predicts that blacks have a higher recidivism rate

Related work and challenges

To mitigate bias associated with protected attributes, many methods have been proposed.



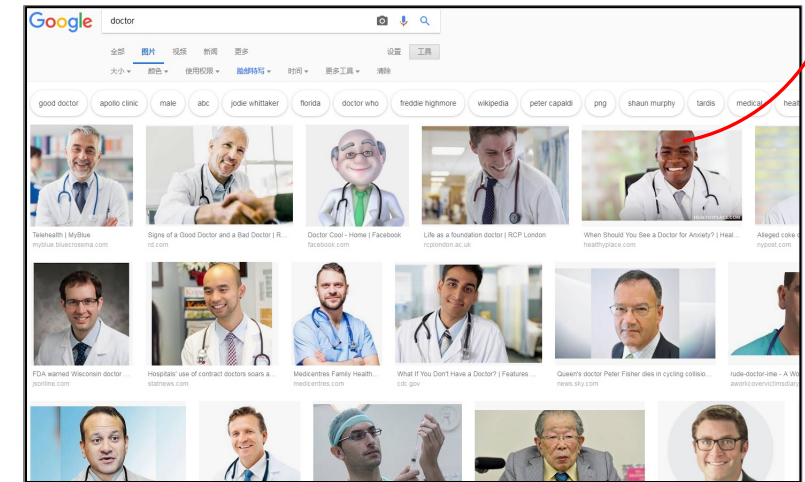
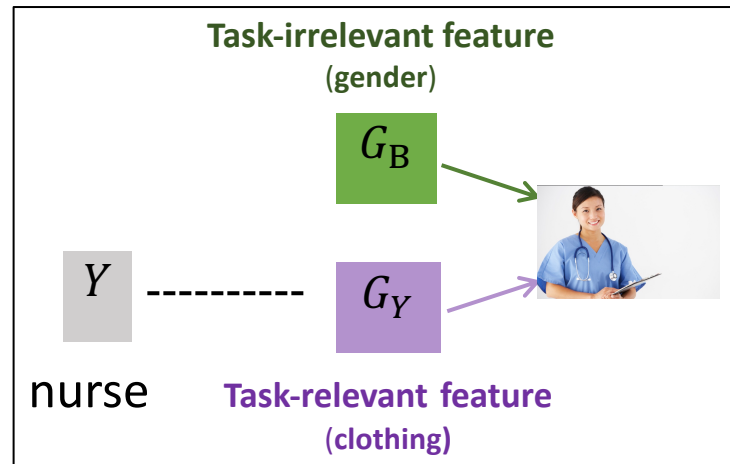
Source of bias: spurious correlations

The algorithm learns spurious correlations from training data :

- ❑ The data samples contain task-irrelevant features (In fairness issues, gender)
- ❑ The task-irrelevant features are strongly related to task-label in the training set
- ❑ The model uses spurious correlations to reduce training loss and turns out model bias

Running example:

Doctor/nurse
Occupation recognition




Retrieved “doctor” images from Google

Eliminate spurious correlations


1 Train a fair model

A trainer with the target of fairness to only use task-relevant features, which makes the learned model fair




2 Get fair results based on a unfair model (more easy-to-use)

A user with the target of fairness to remove the task-irrelevant features of test samples and keep task-irrelevant features, which obtains fair output on unfair model



Part1 : Train a fair model



Adversarial Example-based Data Augmentation for
Eliminating Spurious Correlations in Dataset

Yi Zhang, JitaoSang: Towards Accuracy-Fairness Paradox: Adversarial Example-based Data Augmentation for Visual Debiasing.
ACM Multimedia 2020.

Preliminary

Target variable: Task label to be predicted (occupation: doctor/nurse)

Bias variable: The task-irrelevant attributes that may affect the prediction result (social attributes: gender/skin color)



Female nurse 



Male nurse 

$$\text{bias}(\theta, t) = |P(\hat{t} = t | b = 0, t^* = t) - P(\hat{t} = t | b = 1, t^* = t)|$$

Model Bias: True positive rate (TPR) gap between groups with different **Bias variable** for each **Target variable**

The larger the model bias value, the more unfair the algorithm

Dataset imbalance V.S. Model bias

CelebA Dataset

Target variable: facial attributes

Bias variable: gender


"Arched eyebrows" Recognition



✓



✗



Goal: To balance the data distribution

Up-sampling: assigning different sampling rate to samples, **only guarantees superficial data balance** and can not completely balance

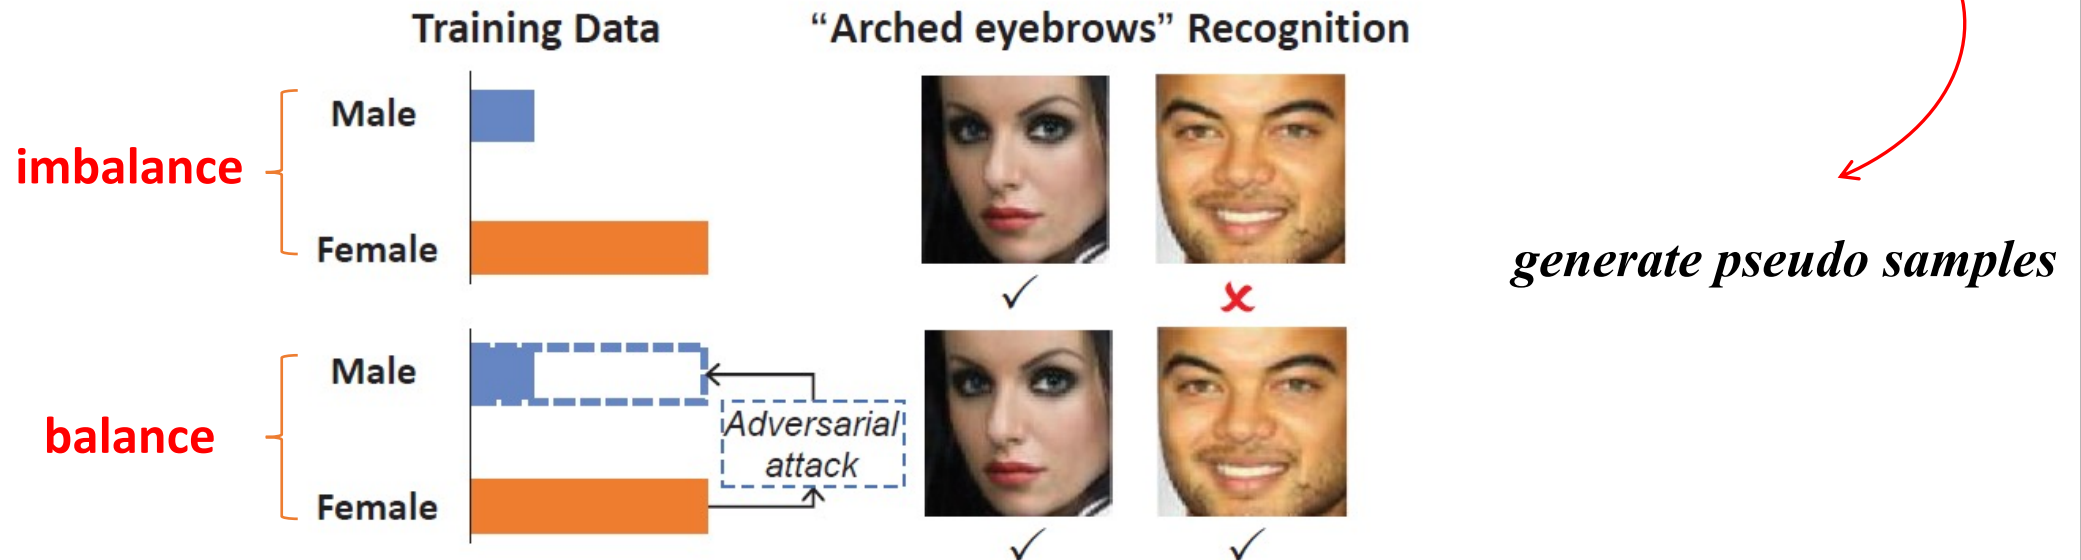
Down-sampling: discarding samples with majority bias variable, **fails to make full use of the training data**

Fairness V.S. Accuracy

Two goals:

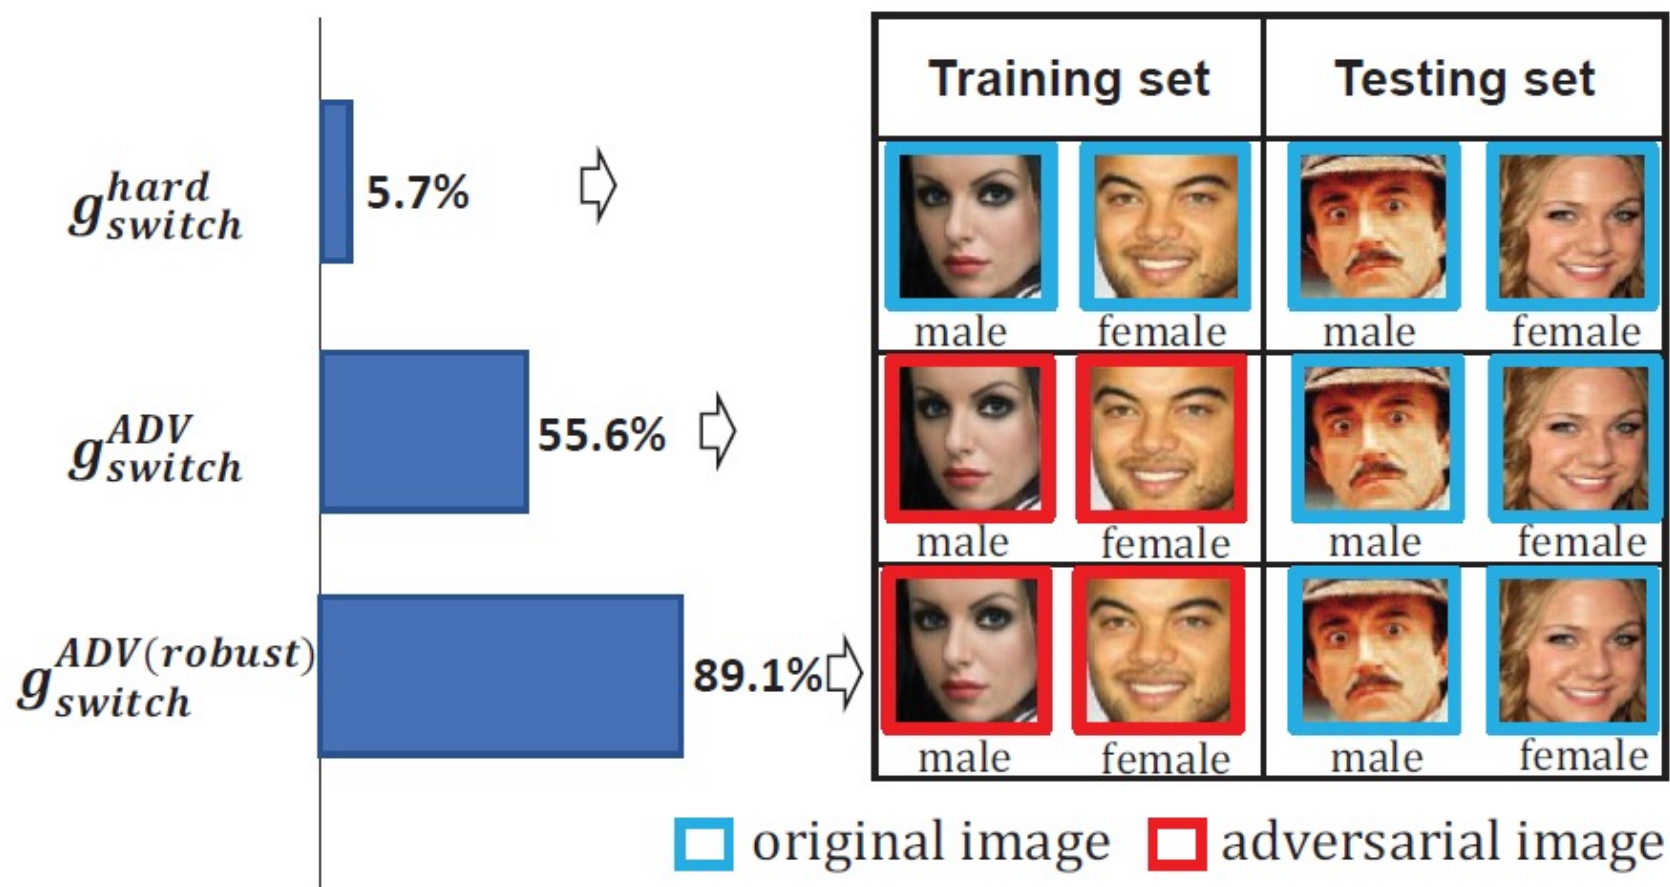
1 **For fairness**, superficial data balance → completely balance

2 **For accuracy**, fails to make full use of the training data → make full use



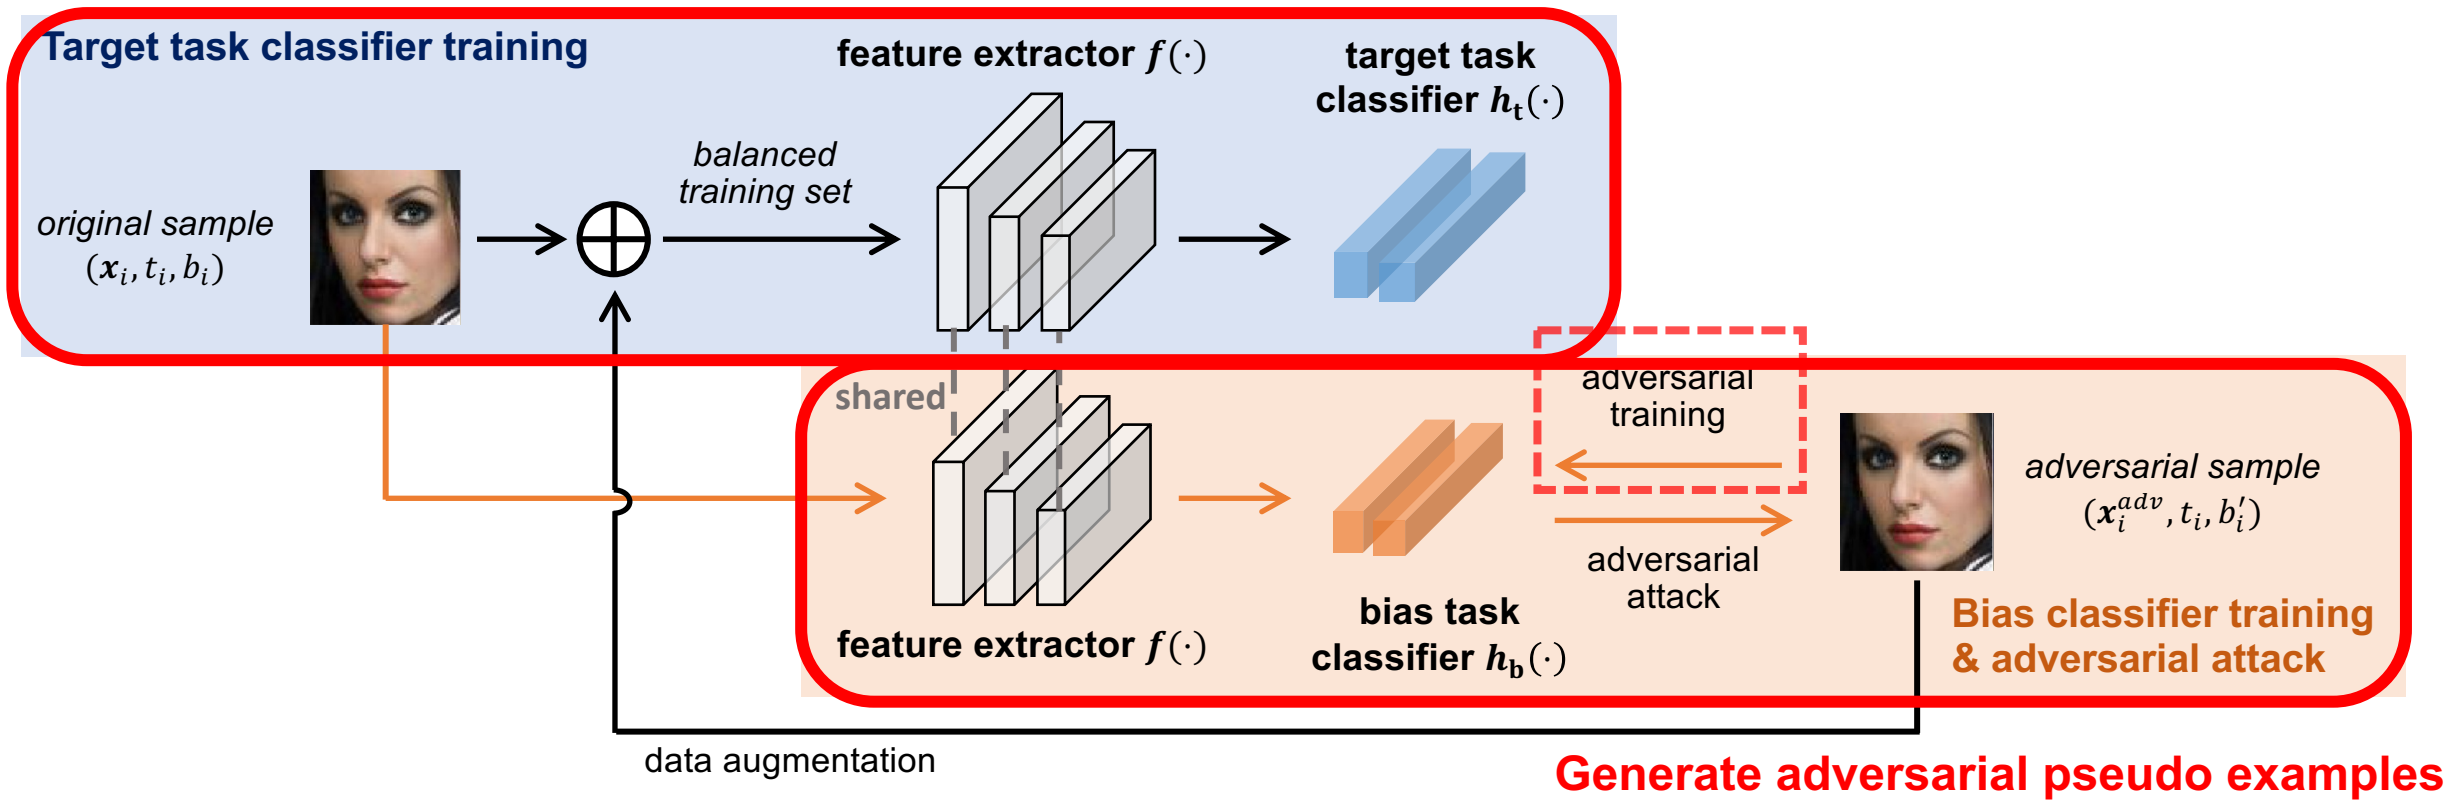
The Potential of Adversarial Example in Balancing

- Adversarial examples **contain useful information about the attack target class** and have potential to **generalize** to original real data



Method

The target task training



Simulated Debiasing Evaluation

□ *target var: digit// bias var: background color*



Methods	bACC (%)	Model bias
Original	55.62	7.84
Under-sampling [7, 30]	—	—
Reweighting [15]	—	—
Adv debiasing [4, 11, 26]	89.93	1.37
CycleGAN [31]	65.23	5.42
AEDA_pre	64.53	5.89
AEDA_online	80.57	3.20
AEDA_robust	91.80	0.53

Accuracy & Model bias in C-MNIST



Confusion matrix for testing subset of
⟨0 ~ 9, red⟩ (left) and ⟨0 ~ 9, brown⟩ (right)

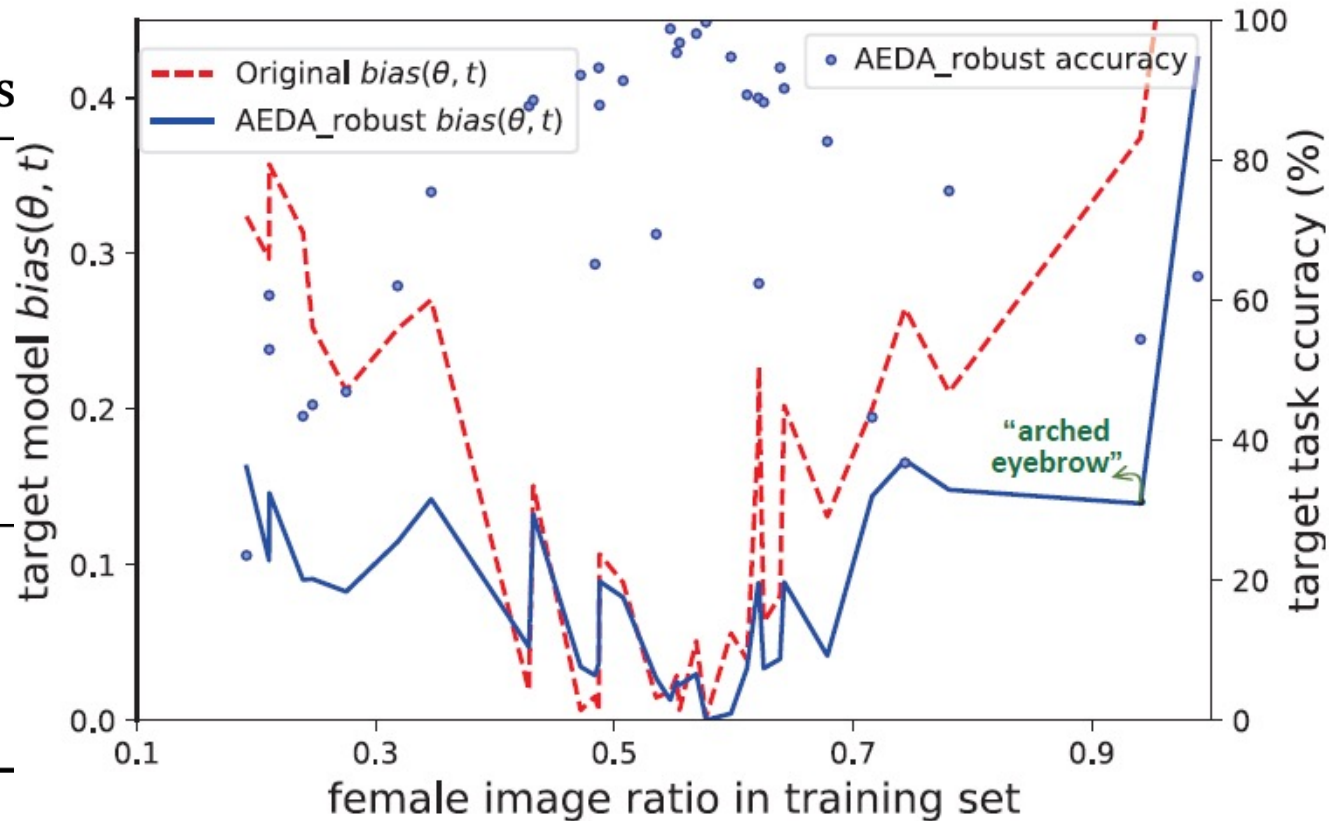
Real-world Debiasing Evaluation

□ *target var*: facial attributes // *bias var*: gender

Methods	bACC (%)	Model bias
Original	73.57	5.48
Under-sampling [7, 30]	66.35	2.35
Reweighting [15]	73.82	4.39
Adv debiasing [4, 11, 26]	72.82	4.23
CycleGAN [31]	73.65	4.75
AEDA_pre	73.68	5.23
AEDA_online	74.03	4.22
AEDA_robust	74.30	3.27

Accuracy & Model bias (gender)

AEDA → Accuracy-compatible Fairness




Accuracy & Model bias (gender) in different facial attributes

Discussion

- *Accuracy-fairness paradox?*
- It is recognized in conventional debiasing attempts that there exists the tradeoff between accuracy and fairness, and the goal is to reduce model bias under the slightly decreased accuracy.
- *Accuracy-compatible Fairness*
- Previous slides: close relation between distribution balance, model bias and accuracy
- Without discarding training data or adding constraints to affect target task learning, the data augmentation provides alternative perspective to simultaneously improve fairness and accuracy.

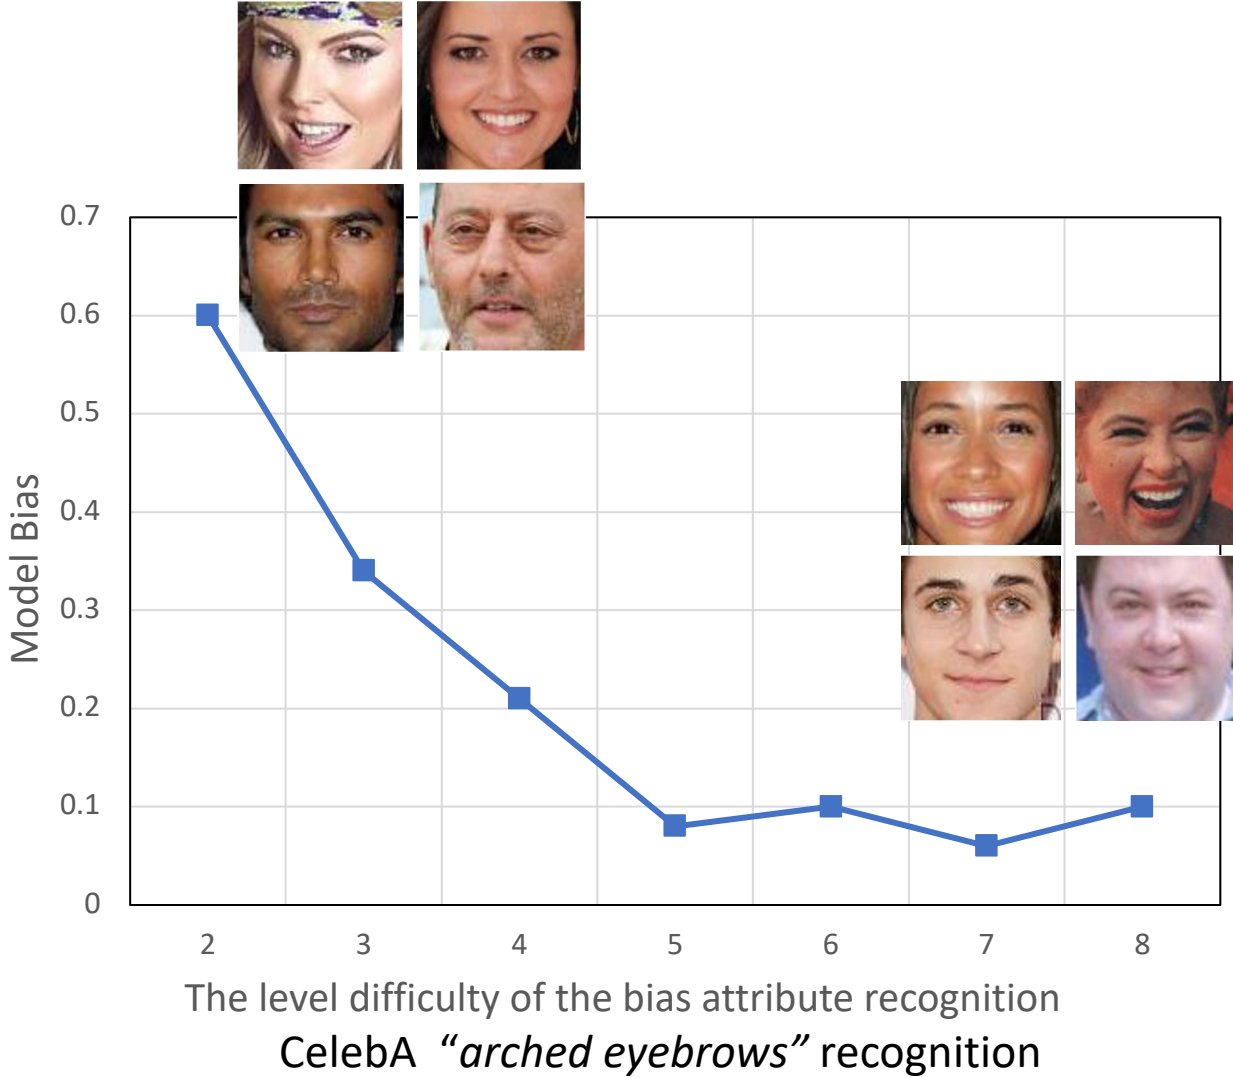
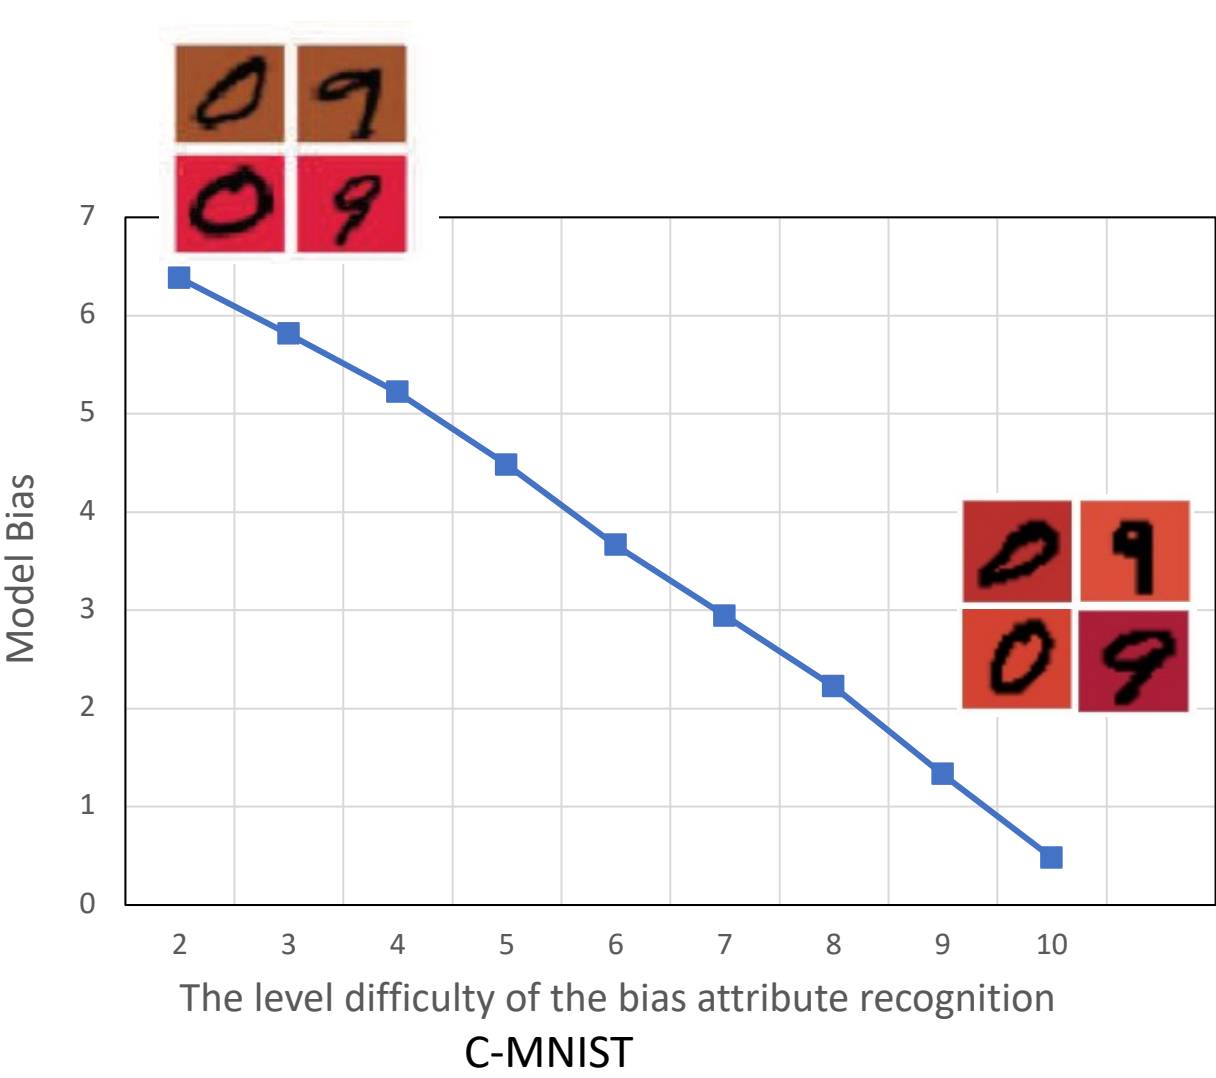
Part2: Get fair result based on unfair model



Post-fairness: Cross-task Adversarial Attack-based Instance Debiasing for Deactivating Spurious Correlations

https://adam-bjtu.org/paper/ZhangYi/Post_fairness_yi/Post_fairness.html

Not all results are based on task-irrelevant features



Regular adversarial attack can't cross task

□ *target var: facial attributes // bias var: gender*

Model Bias	VGG16	VGG19	Resnet18	Resnet32	Resnet50
Original Bias	5.48	5.43	4.93	4.25	5.09
Bias(Instance Debiasing)	4.13	4.26	3.79	3.12	3.47
Bias(Natural Unbiased Instance)	1.98	2.13	1.80	1.96	1.41
$\Delta Bias_{ID}$	1.35	1.17	1.14	1.13	1.62
$\Delta Bias_{NUI}$	3.50	3.30	3.13	2.29	3.68

- Instance Debiasing: using pre-trained gender classifier to remove bias information by adversarial attack.
- Natural Unbiased Instance: searching samples that are difficult for gender classification from the original test set.

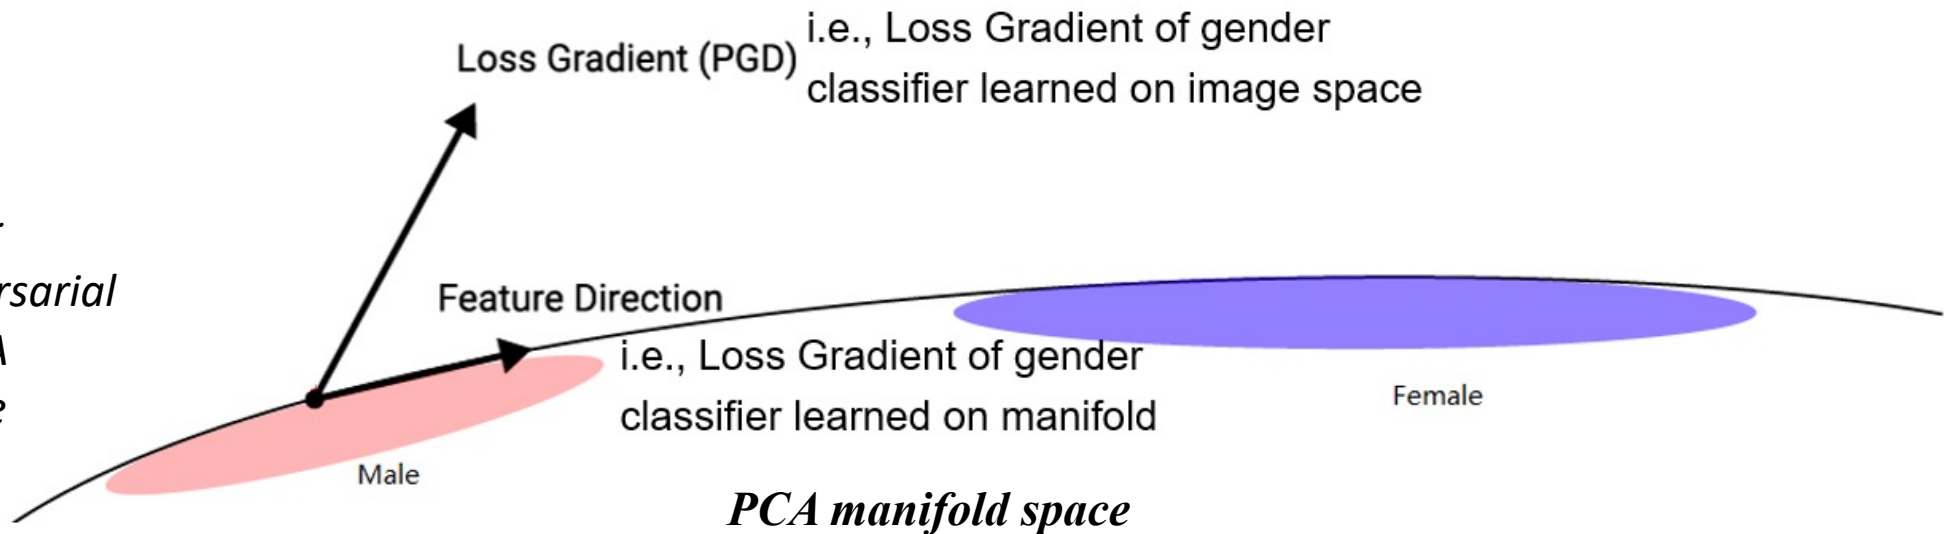


Why does the adversarial attack underperform?

The direction of adversarial attack isn't feature direction

- ❑ Train a gender classifier ***M*** in PCA manifold space.
- ❑ Train a gender classifier ***I*** in image space.
- ❑ In PCA manifold space, the mean Cosine Similarity between “Loss Gradient of ***I***” and “Loss Gradient of ***M***” is **0.36**

Analysis Tool :
Encoding adversarial attack into PCA manifold space

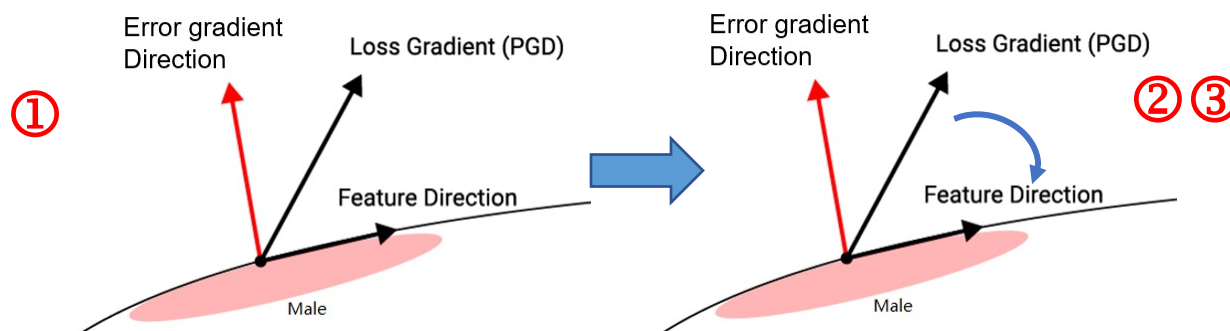


The direction of adversarial attack isn't feature direction

Control the gradient of bias classifier

How to control the gradient of gender classifier to be close to the feature direction?

1. Find the error gradient direction that is not in the feature direction.
2. Only use error gradient direction to generate augmented sample.
3. Use the original sample and the augmented sample to train the bias classifier, and the bias classifier can understand corresponding features of error gradient direction are not the bias features.



Now, the direction of adversarial attack is feature direction

Debiasing Evaluation

Comparison with Post-processing debiasing methods

Dataset	Metric	Vanilla	Ours		Other post-processings		
			Ours [†]	Ours	ROC	EqOdds	CalEqOdds
C-MNIST	Accuracy	55.62	78.90	82.52	N/A	N/A	N/A
	Bias	7.84	1.06	0.74			
CelebA	Accuracy	73.57	73.88	74.32	73.57	68.73	68.63
	Bias	5.48	4.15	2.39	3.52	2.63	5.19

Comparison with Pre and In-processing debiasing methods

Dataset	Metric	Vanilla	Pre-processing		In-processing		Ours
			Re-sampling	Down-sampling	Adv debiasing	AEDA	
C-MNIST	Accuracy	55.62	N/A	N/A	89.93 89.96	91.80 92.32	82.52
	Bias	7.84	N/A	N/A	1.37 0.57	0.53 0.29	0.74
CelebA	Accuracy	73.57	74.19 75.36	64.35 64.67	73.52 74.26	74.30 74.57	74.32
	Bias	5.48	3.83 2.38	2.35 1.92	3.75 2.83	3.27 3.17	2.39

Post-processing as a transitional method

- ***Transformation from basic AI to trustworthy AI***

Long time

- ***AI service provider***

- There may be sufficient obstacles (unbalanced data, time-consuming repeated testing, unpredictable effect, etc.) to be unable to retrain the model


- ***User***

- They want to get the fair predictions about themselves

- ***Third-party***

- They want to get the fair predictions across the population

Post-processing can meet the needs of the three parties, even they can't retrain the model.



Thanks!

Contact me: yi.zhang@bjtu.edu.cn

<http://www.adam-bjtu.org>